

How Spark does it internally?



Ramandeep Kaur
Software Consultant
Knoldus Inc.




Agenda


- Introduction to Apache Spark
- Architecture of Spark cluster
- Spark API: RDD
- Jobs, Stages & Tasks
- DAG
- DAGScheduler
- Shuffling in Spark
- Execution Workflow
- Demo

Apache Spark

- Apache Spark is a fast and general-purpose cluster computing system.
- A lightning-fast cluster computing technology, which is faster than Hadoop MapReduce.

Spark Architecture






SPARK API: RDD


“Fundamental data structure of Apache Spark”

- Resilient
- Distributed
- Dataset
- Lazy evaluated
- Immutable



Operations in Spark

- There are two types of operations:
 - ~ **Transformations:**
 - produces new RDD from the existing RDDs
 - **lazily evaluated**
 - takes RDD as input and produces one or more RDD as output.
 - ~ **Actions:**
 - RDD operations that give non-RDD values.
 - way of sending data from Executer to the driver



TRANSFORMATIONS

map(func)
flatMap(func)
filter()
mapPartitions
mapPartitionWithIndex
union(dataset)
intersection(other-dataset)
distinct()
groupByKey()
reduceByKey(func,
[numTasks])
sortByKey()
join()
coalesce()

ACTIONS

count()
collect()
take(n)
top()
countByValue()
reduce()
fold()
aggregate()
foreach()



Jobs, Stages & Tasks

- **Jobs**

A top-level work item (computation).

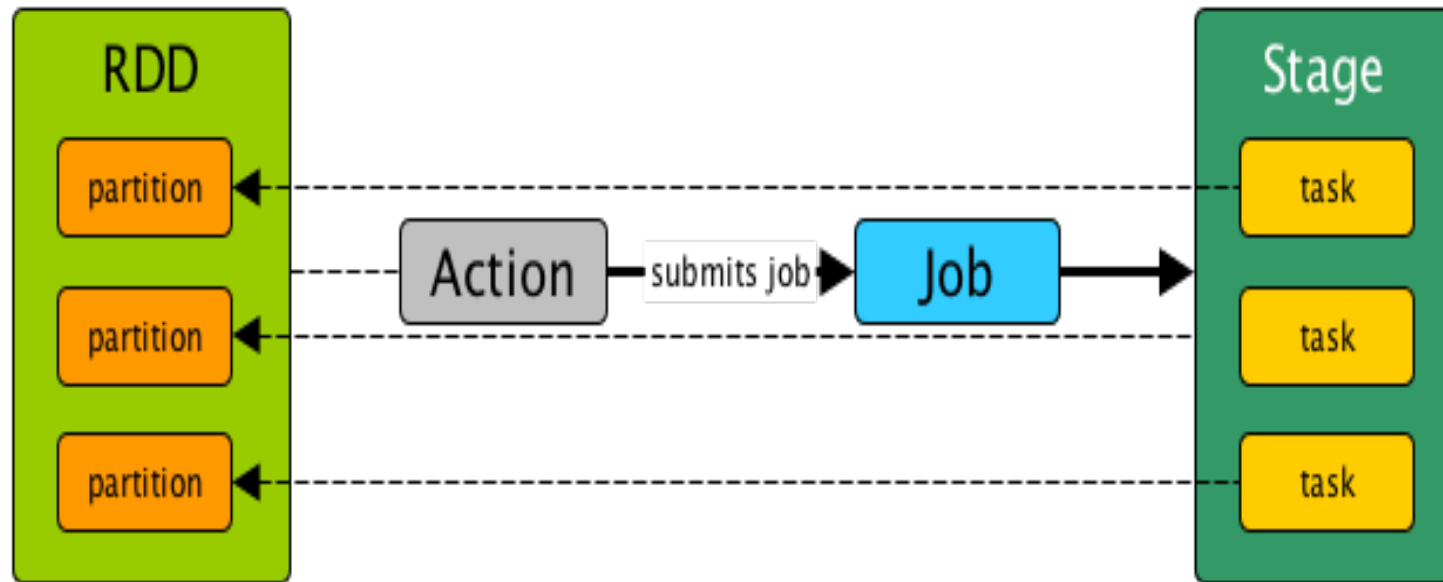
When an action is called the processing gets started and a Job is created which is then submitted to DAGScheduler to be computed.

- **Stages**

Sets of tasks that compute intermediate results in jobs, where each task computes the same function on partitions of the same RDD.

- **Tasks**

A unit of work within a stage, corresponding to one RDD partition.



Transformations


- **Narrow Transformation:** doesn't require data to be shuffled across the partitions.
Eg. Map, filter etc.
- **Wide Transformation:** requires the data to be shuffled among the partitions.
Eg. GroupByKey, reduceByKey etc.

DAG

- **Directed Acyclic Graph**
- Set of Vertices and Edges, where vertices represent the RDDs and the edges represent the Operation to be applied on RDD.

DAG



- DAG is a graph denoting the sequence of operations that are being performed on the target RDD.



DAG

- After an action has been called, **SparkContext** hands over a logical plan to DAGScheduler that it in turn translates to a set of stages that are submitted as TaskSets for execution.

```
(2) MapPartitionsRDD[62] at repartition at <console>:27 □  
  | CoalescedRDD[61] at repartition at <console>:27 □  
  | ShuffledRDD[60] at repartition at <console>:27 □  
+- (8) MapPartitionsRDD[59] at repartition at <console>:27 □  
  | ShuffledRDD[58] at groupBy at <console>:27 □  
+- (8) MapPartitionsRDD[57] at groupBy at <console>:27 □  
  | ParallelCollectionRDD[0] at parallelize at <console>:24 □
```

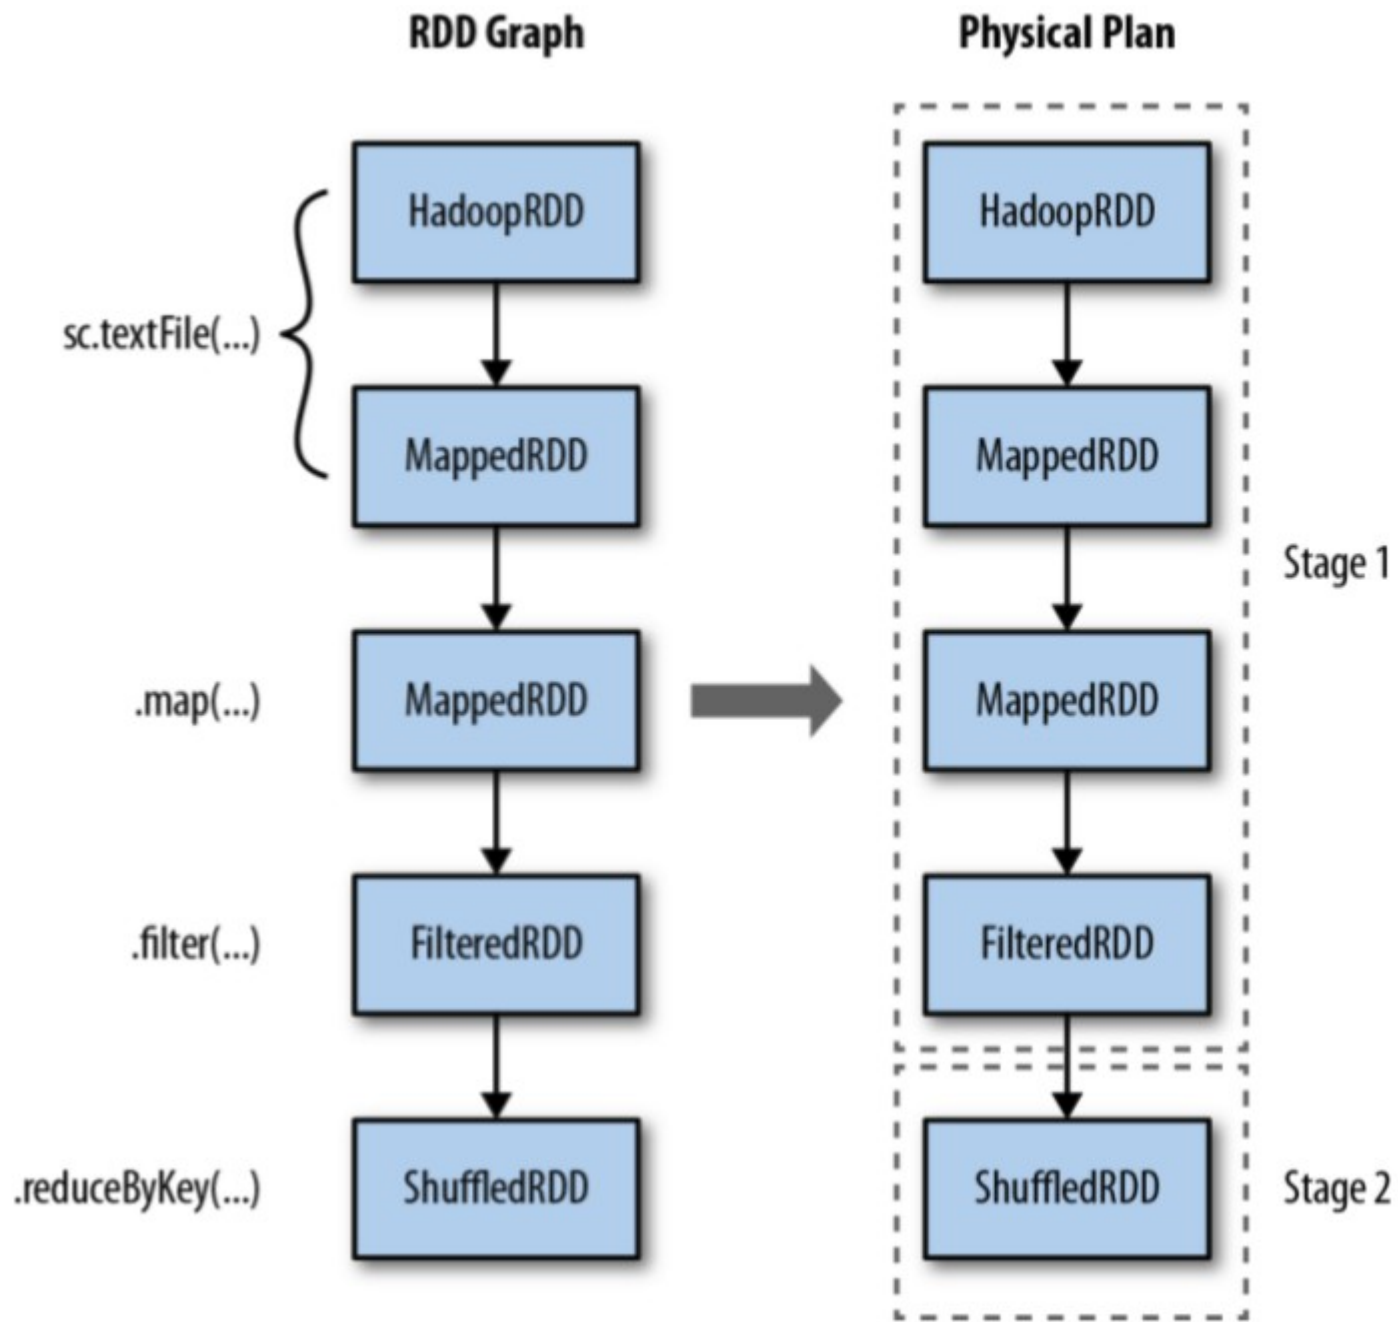



```
scala> val rdd = sc.textFile("/home/knoldus/sparkStreaming.txt")
rdd: org.apache.spark.rdd.RDD[String] = /home/knoldus/sparkStreaming.txt MapPartitionsRDD[1] at textFile at <console>:24

scala> rdd.flatMap(_.split(" ")).filter(_.length > 0)
res0: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[3] at filter at <console>:27

scala> res0.map(x => (x,1))
res1: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[4] at map at <console>:29

scala> res1.reduceByKey((a,b) => a+b)
res2: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[5] at reduceByKey at <console>:31
```



DAGScheduler

- It computes a DAG of stages for each job, keeps track of which RDDs and stage outputs are completed, and finds a minimal schedule to run the job.
- It then submits stages as TaskSets to an underlying TaskScheduler implementation that runs them on the cluster.

- DAGScheduler uses an event queue architecture in which a thread can post **DAGSchedulerEvent** events, e.g. a new job or stage being submitted, that DAGScheduler reads and executes sequentially.

```
private def doOnReceive(event: DAGSchedulerEvent): Unit = event match {  
  case JobSubmitted(jobId, rdd, func, partitions, callSite, listener, properties) =>  
    dagScheduler.handleJobSubmitted(jobId, rdd, func, partitions, callSite, listener, properties)  
  
  case MapStageSubmitted(jobId, dependency, callSite, listener, properties) =>  
    dagScheduler.handleMapStageSubmitted(jobId, dependency, callSite, listener, properties)  
  
  case StageCancelled(stageId, reason) =>  
    dagScheduler.handleStageCancellation(stageId, reason)  
  
  case JobCancelled(jobId, reason) =>  
    dagScheduler.handleJobCancellation(jobId, reason)
```

STAGES

There are two types of Stages:

- **ResultStage**


The final stage in a job that applies a function to one or many partitions of the target RDD to compute the result of an action.

- **ShuffleMapStage**

An intermediate stage in the physical execution DAG that corresponds to a ShuffleDependency.

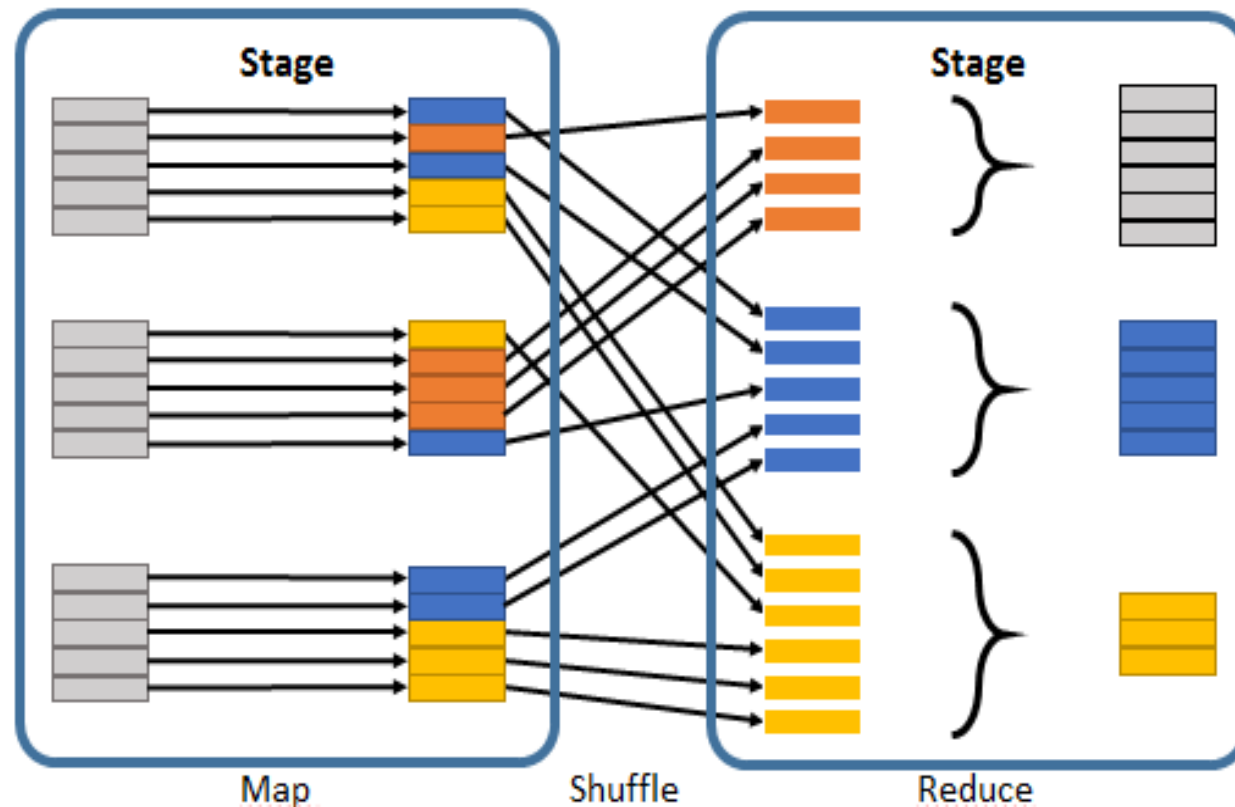
A ShuffleMapStage may contain multiple pipelined operations, e.g. map and filter, before shuffle operation.

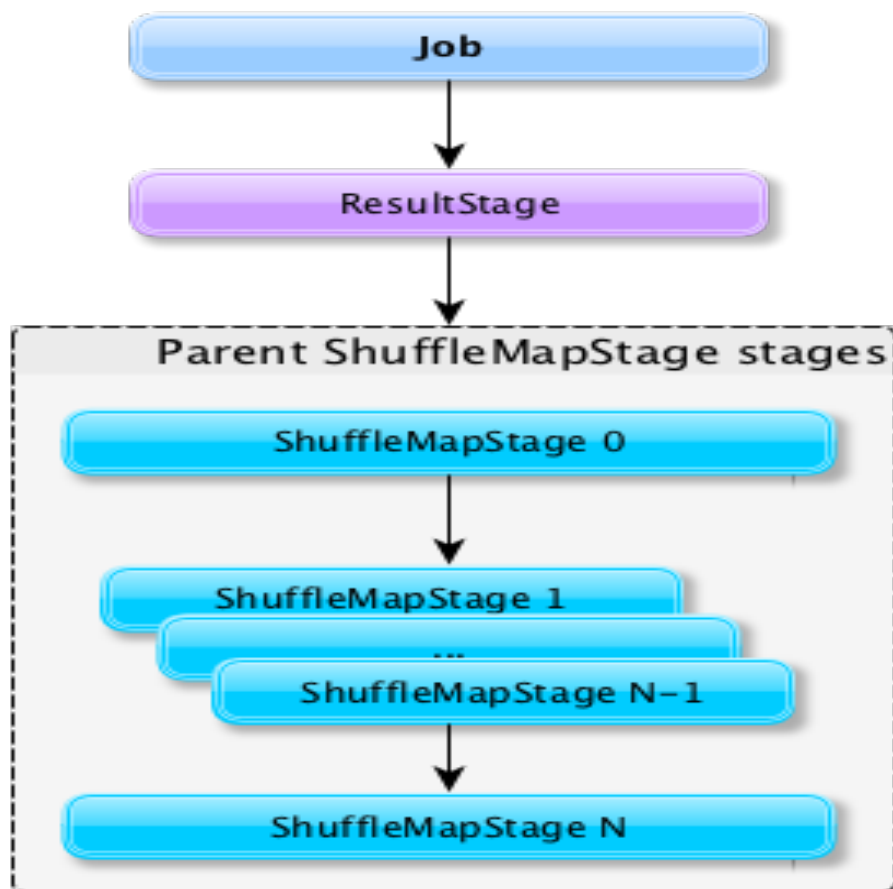
Shuffle Operation



*Process of
redistribution
of data
across
partitions*

Shuffle Map Stage?





There will ALWAYS be only **ONE** ResultStage and can be multiple ShuffleMap stages

Why multiple Stages?

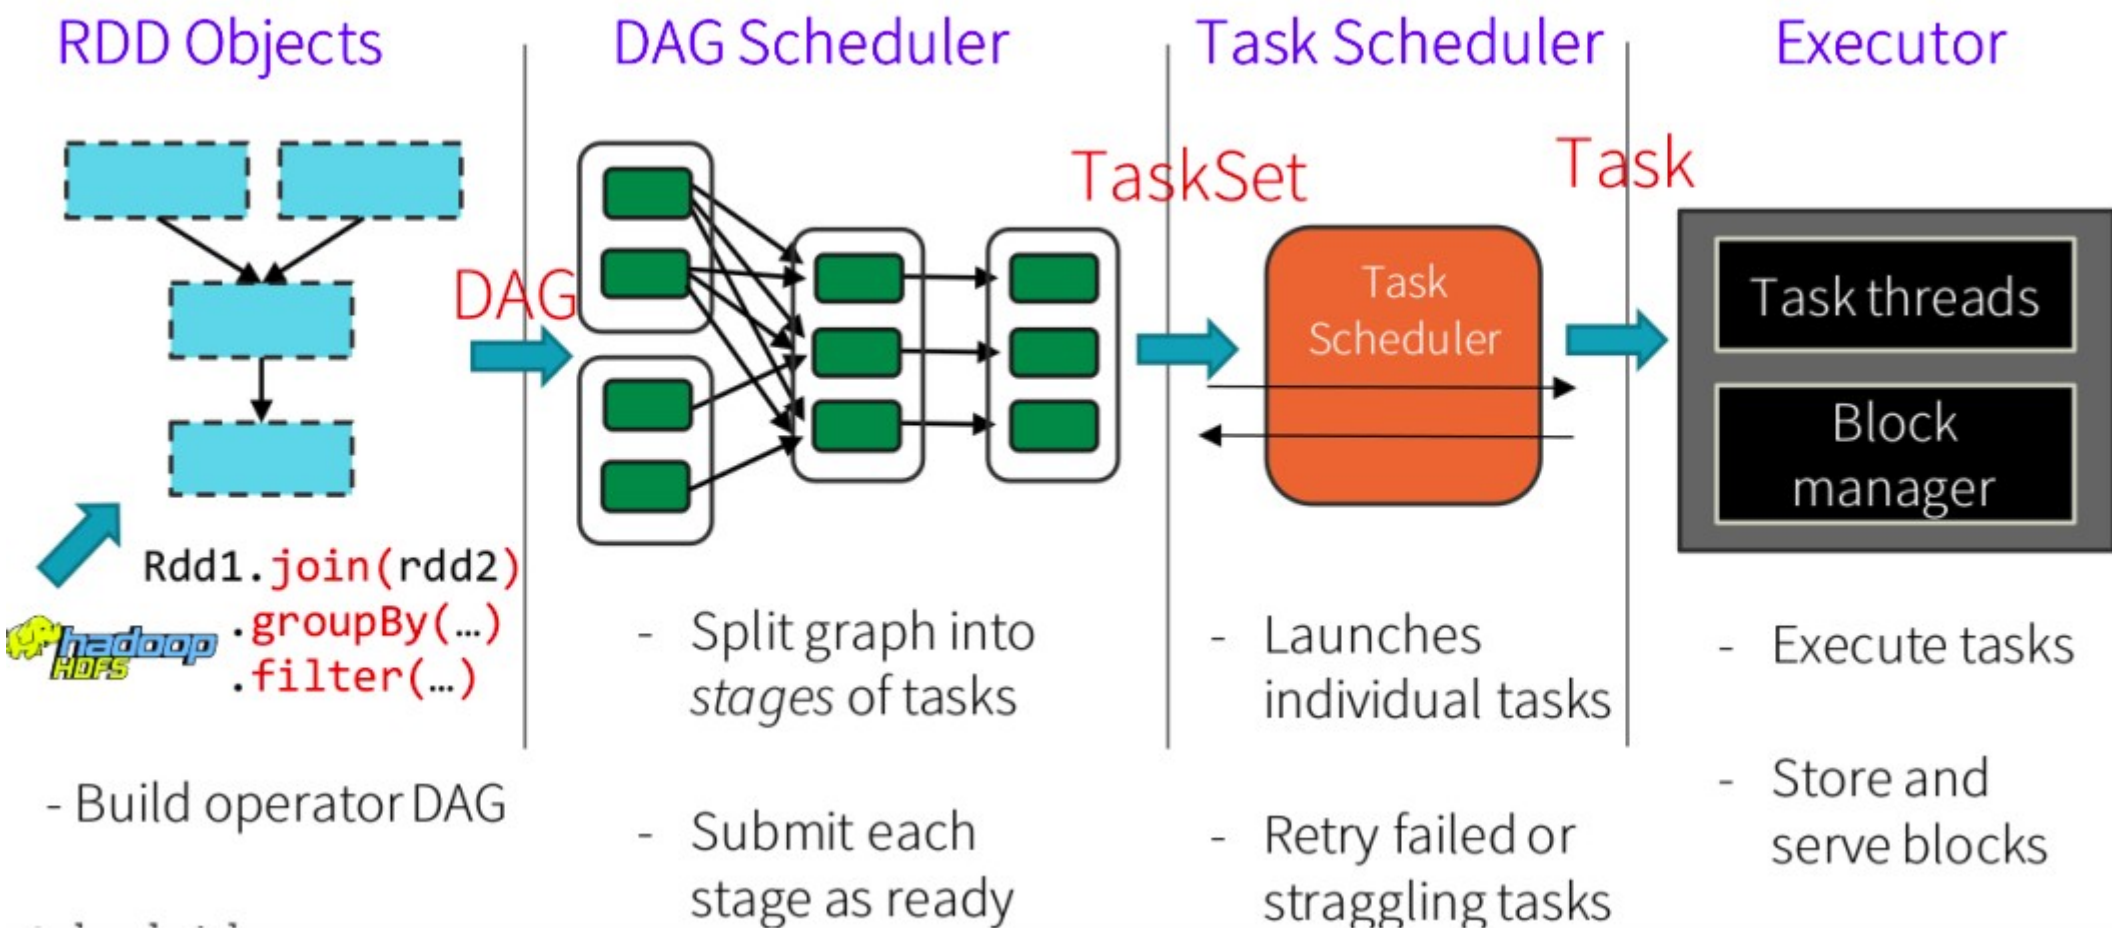
"FAULT TOLERANCE"

- After every shuffle operation, a new stage is created so that whenever data is lost due to shuffle(network I/O) only the previous stage will be calculated for fault tolerance.

Fault Tolerance: HOW?

- To recover from failures, the same stage might need to run multiple times, which are called "attempts".
- If the TaskScheduler reports that a task failed because a map output file from a previous stage was lost, the DAGScheduler resubmits that lost stage.
- This is detected through a **CompletionEvent** with **FetchFailed**, or an **ExecutorLost** event.
- The DAGScheduler will wait a small amount of time to see whether other nodes or tasks fail, then resubmit TaskSets for any lost stage(s) that compute the missing tasks.

Execution Workflow



References

- <https://github.com/apache/spark/blob/master/core/s>
- <https://jaceklaskowski.gitbooks.io/mastering-apach>

Thank you